

Software Dependability Project Report

Improving Code Quality, Security, and Reliability of Apache Commons Codec

Kingstone Showa
Department of Computer Science
University of Salerno
k.showa@studenti.unisa.it

Project: [Apache Commons Codec](#)

January 19, 2025

Introduction

This project focuses on the analysis of Apache Commons Codec, an open-source library developed by the Apache Software Foundation. The library offers robust implementations of essential encoding and decoding algorithms, including Base64, URL encoding and decoding, checksum calculations, and digest algorithms. These functionalities are vital for data transformation and secure communication, making the library a foundational component in numerous software systems.

The primary objective of this project was to evaluate the dependability of the library through a systematic analysis of its source code. This involved identifying potential dependability issues, proposing and implementing improvements, and documenting the results. By analyzing the software's quality, this report provides valuable insights into the library's dependability, highlights areas of concern, and demonstrates how targeted enhancements can bolster its reliability.

Software Quality Analysis

I carried out the software quality analysis of the Apache Commons Codec library using SonarCloud, following these steps,

1. I forked the official Apache Commons Codec repository to my GitHub account to establish a controlled environment for analysis. This step ensured that changes could be made and tested independently of the original repository.
2. I created a new project on SonarCloud and linked to the forked repository. To facilitate continuous integration and automated analysis, the `sonar-project.properties` file was configured, and a CI/CD pipeline was established.
3. SonarCloud analyzed the library's codebase and provided detailed insights into its quality metrics. These included key aspects such as security, reliability, and maintainability.

Issues Categories and Analysis Results

SonarCloud identified issues across three main quality categories: Security, Reliability, and Maintainability.

1. **Security:** No security vulnerabilities were detected, indicating that the library adheres to secure coding practices.
2. **Reliability:** One minor reliability issue was flagged, which demonstrates the robustness of the library. This issue had a negligible impact on the library's overall reliability.
3. **Maintainability:** The majority of the detected issues fell under the maintainability category. These issues were predominantly related to deprecated code snippets and the improper use of switch cases without break commands, which could lead to unintended behavior. SonarCloud identified a total of 1,438 maintainability issues, with varying levels of severity.

Issues Resolved

In resolving these issues, I prioritized issues based on their severity to address the most critical problems first:

1. **Blockers and High-Severity Issues:** All critical issues, such as deprecated code snippets and potentially faulty logic, were resolved. These changes significantly enhanced the code's maintainability and reliability.
2. **Low-Severity Issues:** Many minor issues were resolved to improve overall code quality. However, some were identified as false positives after thorough evaluation.

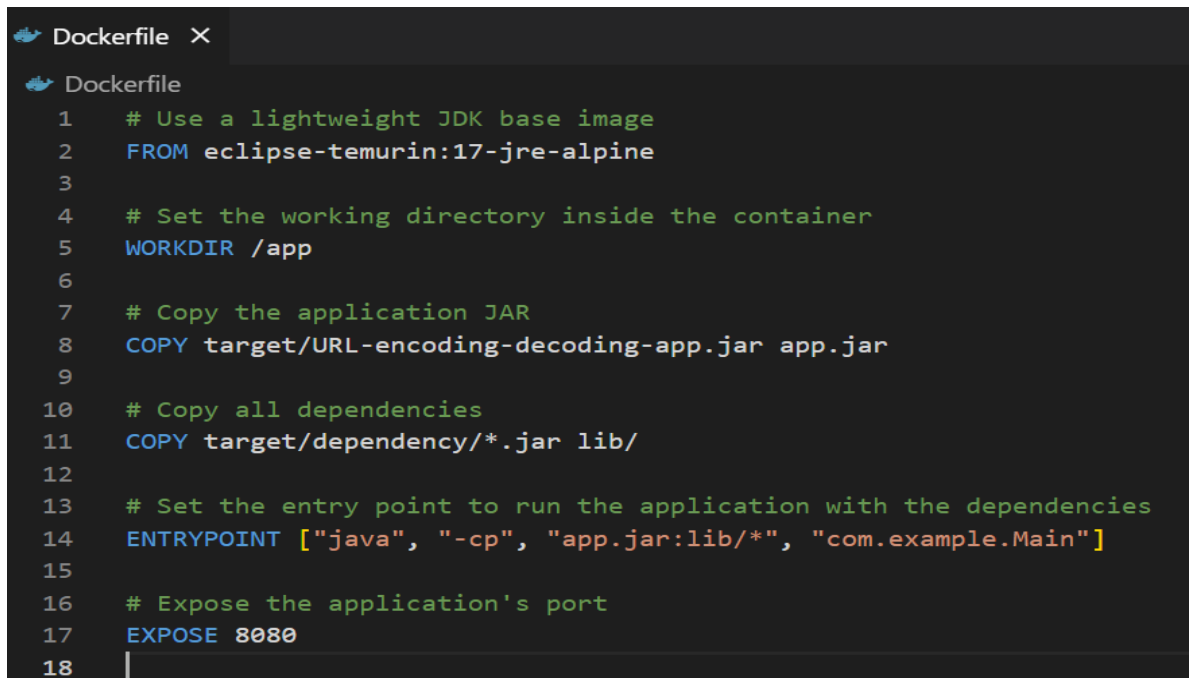
Through these efforts, the total number of maintainability issues was reduced from 1,438 to 1,178, marking a significant improvement in code quality. This reduction highlights the effectiveness of the targeted interventions in enhancing the library's maintainability and overall software quality.

Dockerizing the Text Encoding/Decoding Project

For the second phase of the project, I created a Text Encoding/Decoding Maven application that leverages the updated Apache Commons Codec library. This mini-project provides a simple web-based interface for encoding and decoding text using the library. The application implements Base32 functionality using the Base32 class from the Apache Commons Codec library

Dockerization Process

The application was Dockerized to make it platform-independent (including its dependencies) and to simplify deployment. A Dockerfile was created to define the container's build process and runtime configuration. Below is a Dockerfile and its detailed explanation.



```
1  # Use a lightweight JDK base image
2  FROM eclipse-temurin:17-jre-alpine
3
4  # Set the working directory inside the container
5  WORKDIR /app
6
7  # Copy the application JAR
8  COPY target/URL-encoding-decoding-app.jar app.jar
9
10 # Copy all dependencies
11 COPY target/dependency/*.jar lib/
12
13 # Set the entry point to run the application with the dependencies
14 ENTRYPOINT ["java", "-cp", "app.jar:lib/*", "com.example.Main"]
15
16 # Expose the application's port
17 EXPOSE 8080
18
```

Figure 1: Dockerfile contents

Dockerfile Breakdown

FROM eclipse-temurin:17-jre-alpine

The `eclipse-temurin:17-jre-alpine` image was chosen as the base image. It provides a lightweight Java runtime environment based on JDK 17, optimized for production environments with minimal resource usage.

WORKDIR /app

Sets `/app` as the working directory within the container, where the application and its dependencies will be stored.

COPY target/URL-encoding-decoding-app.jar app.jar

The main application `.jar` file, generated by the Maven build, is copied into the container as `app.jar`.

COPY target/dependency/.jar lib/*

All required `.jar` files (dependencies) from the `target/dependency` folder are copied into a `lib/` directory within the container. This ensures the application can access all required libraries.

ENTRYPOINT ["java", "-cp", "app.jar:lib/", "com.example.Main"]*

The container is configured to execute the application when started. The `java` command is used to run the main class (`com.example.Main`) with the classpath set to include both the application `.jar` and all dependencies in the `lib/` directory.

EXPOSE 8080

Port 8080 is exposed, enabling access to the application from outside the container. This is the default port where the web application listens for incoming requests.

Code Coverage and Mutation Testing

Code Coverage Analysis with JaCoCo

For the code coverage analysis, I used JaCoCo to evaluate how much of the codebase is executed during testing. The initial code coverage for the Apache Commons Codec project was 97.9%, a remarkable result for such a large and complex library. However, gaps were identified, particularly in the *org.apache.commons.codec.cli* package, where the *Digest* class had minimal test coverage at 8.6%.

Element	Coverage	Covered Instructions	Missed Instructions	Total Instructions
commons-codec	97.9 %	180,829	3,891	184,720
src/main/java	97.5 %	48,877	1,232	50,109
org.apache.commons.codec	77.1 %	128	38	166
org.apache.commons.codec	99.1 %	8,178	76	8,254
org.apache.commons.codec.cli	8.6 %	23	243	266
Digest.java	8.6 %	23	243	266
Digest	8.6 %	23	243	266
main(String[])	0.0 %	0	6	6
Digest(String[])	50.0 %	23	23	46
println(String, by)	0.0 %	0	6	6
println(String, by)	0.0 %	0	24	24
run()	0.0 %	0	35	35
run(String, Message)	0.0 %	0	68	68
run(String, Message)	0.0 %	0	28	28

Figure 2: Initial code coverage

To address this, I manually added additional test cases for the *Digest* class, improving its coverage to 18.9%. This increased the overall code coverage of the project to 98.0%.

Element	Coverage	Covered Instructions	Missed Instructions	Total Instructions
commons-codec	98.0 %	180,829	3,747	184,576
src/main/java	97.8 %	48,877	1,088	49,965
org.apache.commons.codec	77.1 %	128	38	166
org.apache.commons.codec	99.1 %	8,178	76	8,254
org.apache.commons.codec.cli	18.9 %	23	99	122
Digest.java	18.9 %	23	99	122
Digest	18.9 %	23	99	122
org.apache.commons.codec	99.0 %	28,633	295	28,928
org.apache.commons.codec	97.2 %	7,448	216	7,664
org.apache.commons.codec	88.8 %	2,524	317	2,841
org.apache.commons.codec	97.6 %	1,943	47	1,990
src/test/java	98.0 %	131,952	2,659	134,611

Figure 3: Improved code coverage

Mutation Testing with PIT

For mutation testing, I used PIT (PITest).The PIT Mutation Testing Report provides detailed metrics on mutation coverage and test strength, as shown in the fig below.

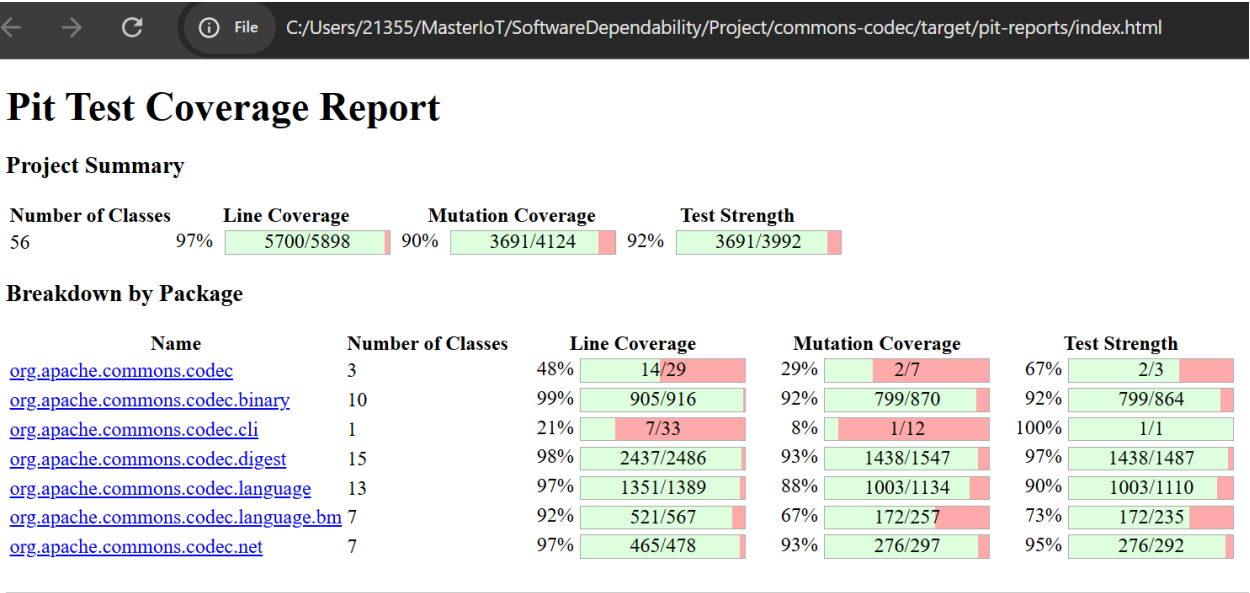


Figure 4: PIT Coverage Report

Metrics from the Report

1. PIT Summary

- Line Coverage: 97% (5,700 out of 5,898 lines are covered by tests).
- Mutation Coverage: 90% (3,691 out of 4,124 mutations were killed by the tests).
- Test Strength: 92% (3,691 out of 3,992 tests detected mutations).

2. Overall Observations

The `org.apache.commons.codec.cli` and `org.apache.commons.codec` packages exhibit weaker metrics compared to other packages. These areas should be prioritized for further testing improvements. The overall mutation coverage of 90% and test strength of 92% demonstrate that the test cases are highly effective at detecting faults, ensuring the robustness of the library.

By leveraging JaCoCo for code coverage and PIT for mutation testing, I increased the overall code coverage from 97.9% to 98.0% by manually adding test cases, particularly for the Digest class, and addressed gaps in mutation detection, achieving a 90% mutation coverage and 92% test strength. This phase highlighted the areas of strength in the project’s test suite while identifying specific packages for further improvement, enhancing the dependability of the Apache Commons Codec library.

Improving Test Coverage

Randoop Test Generation

Randoop was used to generate test cases for the *Digest* class within the *org.apache.commons.codec.cli* package, which initially had minimal test coverage.

Observations After Using Randoop

1. Increase in Digest Class Coverage

The coverage for the *Digest* class increased significantly from 8.6% to 39%, as Randoop generated additional test cases, testing the code paths and functionality of this poorly covered class. This improvement demonstrates the effectiveness of Randoop in targeting under-tested classes.

2. Reduction in Overall Project Coverage

The total project coverage decreased from 98.0% to 96.0%, because Randoop generated a large number of tests, some of which were incomplete, redundant, and less meaningful test cases. These tests did not effectively cover all edge cases, leaving gaps in the coverage.

GitHub Copilot Test Generation

GitHub Copilot, was then used to generate further test cases for the same *Digest* class in the *org.apache.commons.codec.cli* package. Copilot provided more targeted and meaningful tests, improving the depth and quality of the test coverage.

Observations After Using GitHub Copilot

1. Improved Digest Class Coverage

Copilot-generated tests further increased the coverage for the **Digest** class, addressing edge cases and scenarios not covered by the Randoop-generated tests. This increased the *Digest* class coverage significantly beyond the initial improvements made by Randoop.

2. Restoration of Total Project Coverage

The total project coverage increased from 96.0% to 97.8%, nearly matching the initial coverage of 98.0% before Randoop tests were added. Copilot-generated tests were more structured and effective at covering previously missed lines and branches of code, resulting in a higher quality of overall test coverage.

Benchmark Analysis with JHM

The performance of Base32 encoding and decoding was benchmarked using the Java Microbenchmark Harness (JMH). I integrated JMH dependencies into the project using Maven and created a *TextEncodingBenchmark* class with methods for Base32 encoding and decoding benchmarks. The benchmarks were configured to measure throughput (operations per millisecond) over 5 warmup iterations and 5 measurement iterations, each lasting 10 seconds.

```
@BenchmarkMode (Mode.Throughput) // Measures how many operations per unit of time
@OutputTimeUnit (TimeUnit.MILLISECONDS) // Results will be shown in milliseconds
@State (Scope.Thread) // Each thread gets its own state
public class TextEncodingBenchmark {

    private Base32 base32;
    private String inputText;
    private byte[] encoded;
    @Setup (Level.Iteration) // Set up the data once per iteration
    public void setup() {
        base32 = new Base32();
        inputText = "This is a performance test for encoding and decoding using Base32.";
        encoded = base32.encode(inputText.getBytes());
    }

    @Benchmark
    public String testBase32Encoding() {
        return base32.encodeAsString(inputText.getBytes());
    }

    @Benchmark
    public String testBase32Decoding() {return new String(base32.decode(encoded));}
}
```

Figure 5: Base32 Encoding/Decoding benchmarking

Benchmarking Results and Analysis

NOTE: Current JVM experimentally supports Compiler Blackholes, and they are in use. Please exercise extra caution when trusting the results, look into the generated code to check the benchmark still works, and factor in a small probability of new VM bugs. Additionally, while comparisons between different JVMs are already problematic, the performance difference caused by different Blackhole modes can be very significant. Please make sure you use the consistent Blackhole mode for comparisons.

Benchmark	Mode	Cnt	Score	Error	Units
TextEncodingBenchmark.testBase32Decoding	thrpt	5	752.372	47.813	ops/ms
TextEncodingBenchmark.testBase32Encoding	thrpt	5	722.354	55.014	ops/ms

Figure 6: Benchmarking results

Detailed Results

1. Base32 Decoding

- Average Throughput: 752.372 ops/ms
- Variability: ± 47.813 ops/ms (99.9% confidence interval: 704.559 to 800.185 ops/ms)
- Observed Standard Deviation: 12.417 ops/ms

2. Base32 Encoding

- Average Throughput: 722.354 ops/ms
- Variability: ± 55.014 ops/ms (99.9% confidence interval: 667.340 to 777.368 ops/ms)
- Observed Standard Deviation: 14.287 ops/ms

Analysis

The results indicate that Base32 decoding achieves a higher average throughput compared to encoding, suggesting lower computational overhead for decoding operations. Encoding, however, exhibited slightly greater variability, possibly due to the additional complexity in handling input data transformations during the encoding process. Both encoding and decoding demonstrated strong performance, with throughput metrics remaining consistently high. The confidence intervals and standard deviation values confirm that the results are statistically significant and reliable under the test conditions.

Static Security Analysis

1. Find Security Bugs

To analyze security related issues, I installed and run FindBugs on Eclipse. The tool found only one security related issue in *org.apache.commons.codec.digest.PureJavaCrc32Test*.

Issue Detected

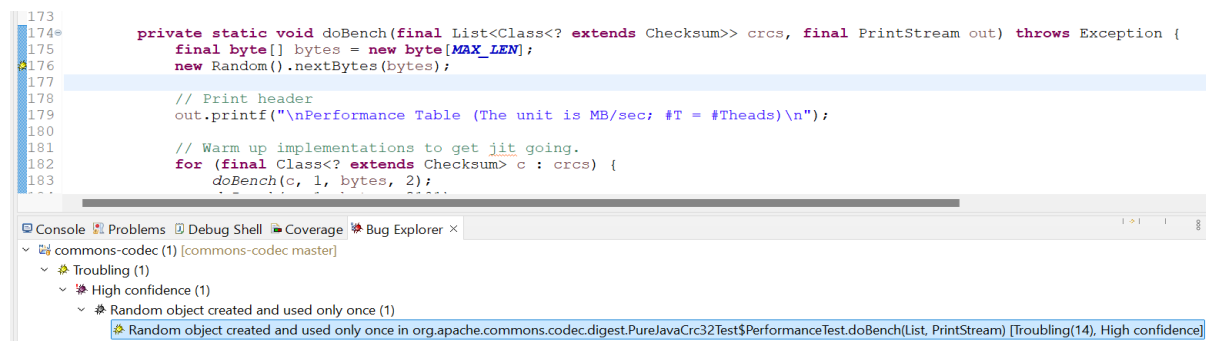


Figure 7: FindBugs execution

Problem details

The code used the *Random* class, which was identified as a potential security risk when generating random values in cryptographic contexts.

Resolution

I replaced the *Random* class with *SecureRandom*, because *SecureRandom* provides cryptographically strong random values suitable for security-sensitive operations, mitigating predictability risks. As a result, the FindBugs tool confirmed that the issue was resolved after making the code change. The project no longer contains high-confidence issues, ensuring a higher level of code reliability and security.

2. OWASP Dependency-Check Analysis

I configured OWASP Dependency-Check and scanned the project for the possible vulnerabilities. The report was generated as follows.

Project: Apache Commons Codec

commons-codec:commons-codec:1.17.2-SNAPSHOT

Scan Information ([show less](#)):

- *dependency-check version*: 8.4.0
- *Report Generated On*: Sat, 18 Jan 2025 12:31:19 +0100
- *Dependencies Scanned*: 19 (19 unique)
- *Vulnerable Dependencies*: 0
- *Vulnerabilities Found*: 0
- *Vulnerabilities Suppressed*: 0
- *NVD CVE Checked*: 2025-01-18T12:29:38
- *NVD CVE Modified*: 2025-01-18T12:00:05
- *VersionCheckOn*: 2025-01-18T12:30:01
- *kev.checked*: 1737199804

Figure 8: OWASP DC Report

The main objective was to analyze project dependencies for known vulnerabilities by cross-referencing them with the National Vulnerability Database (NVD) and other sources.

As a result, no vulnerabilities were identified in the project's dependencies. All libraries are up-to-date and free from reported security issues.

Conclusion

The analysis and improvement of the Apache Commons Codec project have demonstrated the effectiveness of leveraging advanced tools and methodologies to enhance software dependability, security, and quality. By systematically addressing key areas such as maintainability, reliability, and security, the project achieved significant outcomes, summarized as follows,

- | | Software | Quality |
|-----------------------------|---|---------|
| 1. Improved | Using SonarCloud, critical and high-severity issues were identified and resolved, significantly reducing the total number of maintainability issues from 1,438 to 1,178. This improved the library's overall code structure and robustness. | |
| 2. Enhanced Security | Static code analysis with FindBugs identified potential security risks, such as the use of <i>Random</i> , which was replaced with <i>SecureRandom</i> to ensure cryptographically secure operations. | |

Dependency analysis using OWASP Dependency-Check confirmed the absence of vulnerabilities in the project's dependencies, demonstrating the health of the library's ecosystem.

3. Better Test Coverage and Fault Detection

Tools like JaCoCo, PIT mutation testing, Randoop, and GitHub Copilot were instrumental in achieving a high test coverage of 97.8% and a mutation coverage of 90%. These efforts enhanced the library's ability to detect and mitigate faults.

4. Performance

Optimization

Benchmarking with JMH provided insights into the performance of Base32 encoding and decoding, highlighting areas where decoding outperformed encoding. These findings offer opportunities for further optimization.

This project underscores the importance of adopting a multi-faceted approach to software analysis and improvement, combining static analysis, dynamic testing, dependency scanning, and benchmarking. The enhancements implemented have not only improved the current state of the Apache Commons Codec library but also established a strong foundation for its future development and maintenance.

References

1. **Apache Software Foundation** - Apache Commons Codec Documentation <https://commons.apache.org/proper/commons-codec/>
2. **SonarCloud Documentation** - Continuous Integration and Code Quality <https://sonarcloud.io>
3. **JaCoCo Documentation** - Java Code Coverage Library <https://www.jacoco.org/jacoco/>
4. **OWASP Dependency-Check** - Vulnerability Scanning Tool <https://owasp.org/www-project-dependency-check/>
5. **PIT Mutation Testing** - Fault Detection in Java Applications <https://pitest.org/>
6. **Java Microbenchmark Harness (JMH)** - Benchmarking Framework <https://openjdk.org/projects/code-tools/jmh/>
7. **FindBugs** - Static Analysis Tool for Java <http://findbugs.sourceforge.net/>
8. **GitHub Copilot Documentation** - AI-Powered Code Assistant <https://github.com/features/copilot>
9. **Randoop** - Test Case Generation Tool <https://github.com/randoop/randoop/releases/tag/v4.3.2>